

**Face recognition** is the general task of identifying and verifying people from photographs of their face. But nowadays and with the high propagation of the **COVID-19**, this task became more and more difficult due to the **masks** weared by the people. In this project, we will discover how to develop a **face recognition system** using **FaceNet** and an SVM and KNN classifiers to

## Import the required libraries

```
In [1]: from IPython.display import clear_output
import numpy as np
from numpy import expand_dims
from numpy import asarray
from PIL import Image

from sklearn.preprocessing import Normalizer
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix

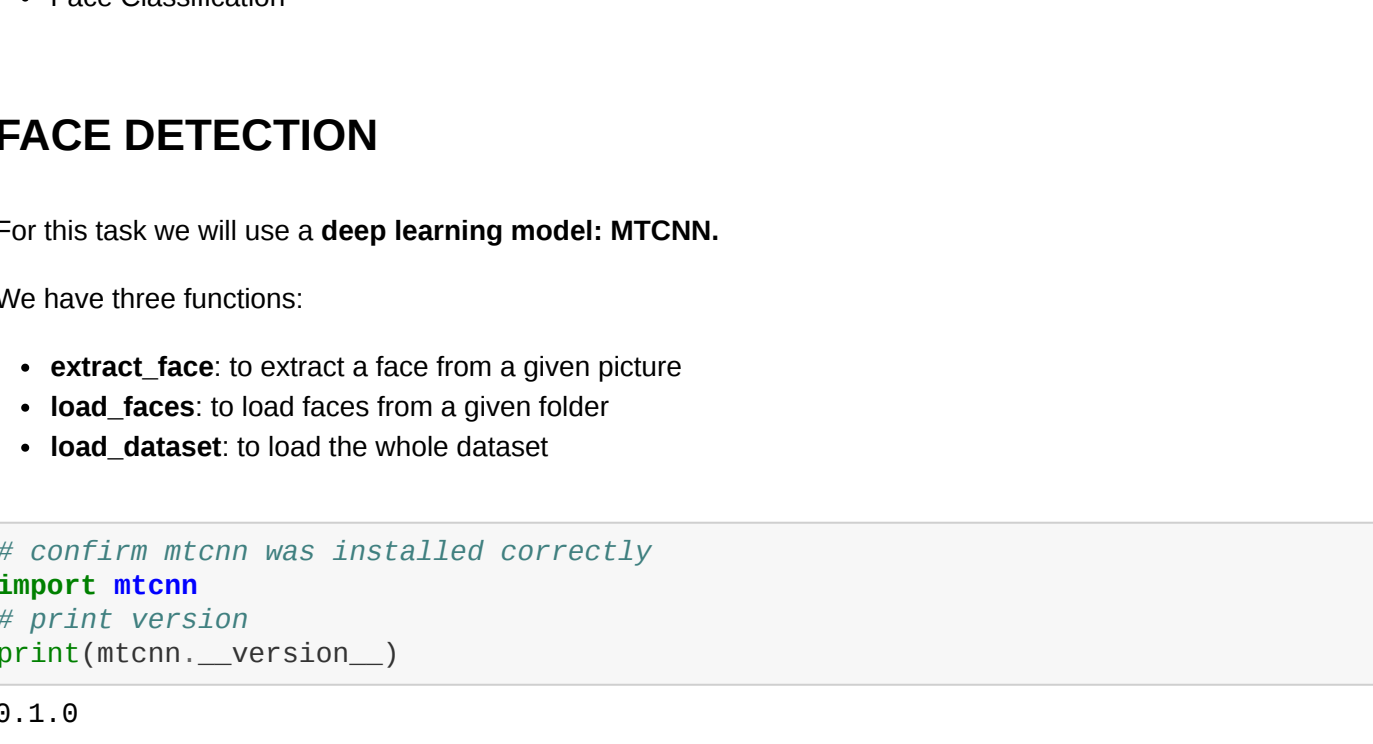
from sklearn import svm
from sklearn.neighbors import KNeighborsClassifier

from matplotlib import pyplot
import matplotlib.pyplot as plt
from mtcnn.mtcnn import MTCNN
from os import listdir

# example of loading the keras facenet model
from keras.models import load_model
```

## DataSet

Our data set is composed of two folders: **Train** and **Test**. Each one contains **16 folders of famous people and some friends**. For the train, each person has **15 pictures** in average while for the test we have **5 pictures** in average. We took the **Train Dataset** and we passed it through an algorithm that takes a **non masked person as an input** and gives us as an output the **same picture but with mask**.



The work is divided into three parts:

- Face Detection
- Face Embedding
- Face Classification

## FACE DETECTION

For this task we will use a **deep learning model: MTCNN**.

We have three functions:

- **extract\_face**: to extract a face from a given picture
- **load\_faces**: to load faces from a given folder
- **load\_dataset**: to load the whole dataset

```
In [3]: # confirm mtcnn was installed correctly
import mtcnn
# print version
print(mtcnn.__version__)

0.1.0

In [4]: # extract a single face from a given photograph
def extract_face(filename, required_size=(160, 160)):
    # load image from file
    image = Image.open(filename)
    # convert to RGB, if needed
    image = image.convert('RGB')
    # convert to array
    pixels = asarray(image)
    # create the detector, using default weights
    detector = MTCNN()
    # detect faces in the image
    results = detector.detect_faces(pixels)
    # extract the bounding box from the first face
    x1, y1, width, height = results[0]['box']
    # bug fix
    x2, y2 = x1 + width, y1 + height
    # extract the face
    face = pixels[y1:y2, x1:x2]
    # resize pixels to the model size
    image = Image.fromarray(face)
    image = image.resize(required_size)
    face_array = asarray(image)
    return face_array

In [5]: # load images and extract faces for all images in a directory
def load_faces(directory):
    faces = list()
    # enumerate files
    for filename in listdir(directory):
        # path
        path = directory + filename
        # get face
        face = extract_face(path)
        # store
        faces.append(face)
    return faces

In [6]: # load a dataset that contains one subdir for each class that in turn contains images
def load_dataset(directory):
    X, y = list(), list()
    # enumerate classes on per class
    for subdir in listdir(directory):
        # path
        path = directory + subdir + '/'
        # skip any files that might be in the dir
        #if not isdir(path):
            #continue
        # load all faces in the subdirectory
        faces = load_faces(path)
        # create labels
        labels = [subdir for _ in range(len(faces))]
        # summarize progress
        print('>loaded %d examples for class: %s' % (len(faces), subdir))
        # store
        X.extend(faces)
        y.extend(labels)
    return asarray(X), asarray(y)

In [7]: # load train dataset
trainX, trainy = load_dataset('train/')
print(trainX.shape, trainy.shape)
clear_output()

In [8]: # load test dataset
testX, testy = load_dataset('val/')
print(testX.shape, testy.shape)
clear_output()

In [9]: # save arrays to one file in compressed format
np.savez_compressed('masked-faces-dataset.npz', trainX, trainy, testX, testy)

In [10]: # load the face dataset
data = np.load('masked-faces-dataset.npz')
trainX, trainy, testX, testy = data['arr_0'], data['arr_1'], data['arr_2'], data['arr_3']
print('Loaded: ', trainX.shape, trainy.shape, testX.shape, testy.shape)

Loaded: (427, 160, 160, 3) (427,) (105, 160, 160, 3) (105,)
```

As mentioned in the output of the last cell: we have **427 pictures for the train** and **105 for the test**.

## Face Embeddings

In this part, We will use **Face Net**: it is a system that, given a picture of a face, will extract high-quality features from the face and predict a 128 element vector representation these features, called a face embedding. These **face embeddings** will then be used as the basis for training classifier systems on standard face recognition benchmark datasets

```
In [11]: # load the model
model = load_model('facenet_keras.h5')

WARNING:tensorflow:No training configuration found in the save file, so the model was 'not' compiled. Compile it manually.

In [12]: # get the face embedding for one face
def get_embedding(model, face_pixels):
    # scale pixel values
    face_pixels = face_pixels.astype('float32')
    # standardize pixel values across channels (global)
    mean, std = face_pixels.mean(), face_pixels.std()
    face_pixels = (face_pixels - mean) / std
    # transform face into one sample
    samples = expand_dims(face_pixels, axis=0)
    # make prediction to get embedding
    yhat = model.predict(samples)
    return yhat[0]

In [13]: # convert each face in the train set to an embedding
newTrainX = list()
for face_pixels in trainX:
    embedding = get_embedding(model, face_pixels)
    newTrainX.append(embedding)
newTrainX = asarray(newTrainX)
print(newTrainX.shape)

(427, 128)

In [14]: # convert each face in the test set to an embedding
newTestX = list()
for face_pixels in testX:
    embedding = get_embedding(model, face_pixels)
    newTestX.append(embedding)
newTestX = asarray(newTestX)
print(newTestX.shape)

(105, 128)

In [15]: # save arrays to one file in compressed format
np.savez_compressed('masked-faces-embeddings.npz', newTrainX, trainy, newTestX, testy)

In [16]: # load dataset
data = np.load('masked-faces-embeddings.npz')
trainX, trainy, testX, testy = data['arr_0'], data['arr_1'], data['arr_2'], data['arr_3']
print('Dataset: train=%d, test=%d' % (trainX.shape[0], testX.shape[0]))

Dataset: train=427, test=105

In [17]: # normalize input vectors
in_encoder = Normalizer(norm='l2')
trainX = in_encoder.transform(trainX)
testX = in_encoder.transform(testX)

In [18]: # label encode targets
out_encoder = LabelEncoder()
out_encoder.fit(trainy)
trainy = out_encoder.transform(trainy)
testy = out_encoder.transform(testy)
```

## Face Classification

### SVC

```
In [19]: # fit model
model = svm.SVC(kernel='linear')
model.fit(trainX, trainy)

Out[19]: SVC(C=1.0, cache_size=200, class_weight=None, coef=0.0,
decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
kernel='linear', max_iter=1, probability=False, random_state=None,
shrinking=True, tol=0.001, verbose=False)

In [20]: # predict
yhat_train = model.predict(trainX)
yhat_test = model.predict(testX)
```

#### Evaluation for the Train

```
In [21]: print ('Accuracy_train:', accuracy_score(trainy, yhat_train))
print ('F1 score_train:', f1_score(trainy, yhat_train, average='weighted'))
print ('Recall_train:', recall_score(trainy, yhat_train,
average='weighted'))
print ('Precision_train:', precision_score(trainy, yhat_train,
average='weighted'))
print ('\n classification report_train:\n', classification_report(trainy, yhat_train))
print ('\n confusion matrix_train:\n', confusion_matrix(trainy, yhat_train))

Accuracy_train: 0.990632318501171
F1 score_train: 0.990629386649055
Recall_train: 0.990632318501171
Precision_train: 0.9913785864100476

classification report_train:
      precision    recall  f1-score   support

   0      0.99      1.00      1.00      122
   1      1.00      1.00      1.00      125
   2      1.00      0.90      0.95       10
   3      1.00      0.89      0.94        0
   4      1.00      1.00      1.00       13
   5      0.90      1.00      0.95        9
   6      1.00      1.00      1.00        0
   7      1.00      0.93      0.97       15
   8      1.00      1.00      1.00       17
   9      1.00      1.00      1.00       21
  10      1.00      1.00      1.00       11
  11      1.00      1.00      1.00       15
  12      1.00      0.95      0.98       21
  13      0.89      1.00      0.94       17
  14      1.00      1.00      1.00       13

   accuracy      0.99      0.98      0.99      427
  macro avg      0.99      0.98      0.98      427
 weighted avg      0.99      0.99      0.99      427

confusion matrix_train:
[[122  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0 125  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  9  0  0  0  0  0  0  0  0  0  0  1  0]
 [  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  13  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  9  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  9  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0 14  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0 17  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0  0 21  0  0  0  0  0]
 [  0  0  0  0  0  0  0  0  0 11  0  0  0  0]
 [  0  0  0  0  0  1  0  0  0  0  0 15  0  0]
 [  0  0  0  0  0  0  1  0  0  0  0  0 20  0  0]
 [  0  0  0  0  0  0  0  0  0  0  0  0  0 17  0]
 [  0  0  0  0  0  0  0  0  0  0  0  0  0  0 13]]
```

#### Evaluation for the Test

```
In [22]: print ('Accuracy_test:', accuracy_score(testy, yhat_test))
print ('F1 score_test:', f1_score(testy, yhat_test, average='weighted'))
print ('Recall_test:', recall_score(testy, yhat_test,
average='weighted'))
print ('Precision_test:', precision_score(testy, yhat_test,
average='weighted'))
print ('\n classification report_test:\n', classification_report(testy, yhat_test))
print ('\n confusion matrix_test:\n', confusion_matrix(testy, yhat_test))

Accuracy_test: 0.8285714285714286
F1 score_test: 0.8159095864377997
Recall_test: 0.8285714285714286
Precision_test: 0.8651251526251525

classification report_test:
      precision    recall  f1-score   support

   0      0.77      1.00      0.87       10
   1      0.94      1.00      0.97       15
   2      1.00      1.00      1.00        5
   3      1.00      1.00      1.00        5
   4      0.83      1.00      0.91        5
   5      1.00      0.40      0.57        5
   6      1.00      1.00      1.00        5
   7      0.50      0.78      0.67        9
   8      0.83      1.00      0.91        5
   9      1.00      1.00      1.00        5
  10      1.00      0.44      0.62        9
  11      1.00      1.00      1.00        7
  12      1.00      1.00      1.00        5
  13      0.50      0.86      0.63        7
  14      0.75      0.58      0.50        8

   accuracy      0.85      0.84      0.83      105
  macro avg      0.85      0.84      0.83      105
 weighted avg      0.87      0.83      0.82      105

confusion matrix_test:
[[10  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0 15  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  5  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  5  0  0  0  0  0  0  0  0  0  0  0  0]
 [  2  0  0  0  2  0  0  0  0  0  0  1  1  0  2]
 [  0  0  0  0  5  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  7  1  0  0  0  0  0  0  0  1]
 [  0  0  0  0  0  0  5  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0  5  0  0  0  0  0  0  0]
 [  0  1  0  0  1  0  1  0  4  0  0  2  0  0  0]
 [  0  0  0  0  0  0  1  0  0  0  5  0  1  0  0]
 [  0  0  0  0  0  0  0  0  0  0  0  0  5  0  0]
 [  1  0  0  0  0  0  0  0  0  0  0  0  0  0  6]
 [  0  0  0  0  0  0  3  0  0  0  0  0  2  3  3]]
```

### KNN

```
In [23]: knn = KNeighborsClassifier(n_neighbors=1, metric='euclidean')
knn.fit(trainX, trainy)

Out[23]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='euclidean',
metric_params=None, n_jobs=None, n_neighbors=1, p=2,
weights='uniform')

In [24]: from sklearn.metrics import accuracy_score
# predict
yhat_train = knn.predict(trainX)
yhat_test1 = knn.predict(testX)
```

#### Evaluation for the Train

```
In [25]: print ('Accuracy_train:', accuracy_score(trainy, yhat_train))
print ('F1 score_train:', f1_score(trainy, yhat_train, average='weighted'))
print ('Recall_train:', recall_score(trainy, yhat_train,
average='weighted'))
print ('Precision_train:', precision_score(trainy, yhat_train,
average='weighted'))
print ('\n classification report_train:\n', classification_report(trainy, yhat_train))
print ('\n confusion matrix_train:\n', confusion_matrix(trainy, yhat_train))

Accuracy_train: 1.0
F1 score_train: 0.990629386649055
Recall_train: 1.0
Precision_train: 1.0

classification report_train:
      precision    recall  f1-score   support

   0      1.00      1.00      1.00      122
   1      1.00      1.00      1.00      125
   2      1.00      1.00      1.00       10
   3      1.00      1.00      1.00        9
   4      1.00      1.00      1.00       13
   5      1.00      1.00      1.00        0
   6      1.00      1.00      1.00        0
   7      1.00      1.00      1.00       15
   8      1.00      1.00      1.00       17
   9      1.00      1.00      1.00       21
  10      1.00      1.00      1.00       11
  11      1.00      1.00      1.00       15
  12      1.00      1.00      1.00       21
  13      1.00      1.00      1.00       17
  14      1.00      1.00      1.00       13

   accuracy      1.00      1.00      1.00      427
  macro avg      1.00      1.00      1.00      427
 weighted avg      1.00      1.00      1.00      427

confusion matrix_train:
[[122  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0 125  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0 10  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  9  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0 13  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  9  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0  9  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0 14  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0  0 17  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0  0  0 17  0  0  0  0  0]
 [  0  0  0  0  0  0  0  0  0  0 11  0  0  0  0]
 [  0  0  0  0  0  0  0  0  0  0  0 15  0  0  0]
 [  0  0  0  0  0  0  0  0  0  0  0  0 21  0  0]
 [  0  0  0  0  0  0  0  0  0  0  0  0  0 17  0]
 [  0  0  0  0  0  0  0  0  0  0  0  0  0  0 13]]
```

#### Evaluation for the Test

```
In [26]: print ('Accuracy_test:', accuracy_score(testy, yhat_test1))
print ('F1 score_test:', f1_score(testy, yhat_test1, average='weighted'))
print ('Recall_test:', recall_score(testy, yhat_test1,
average='weighted'))
print ('Precision_test:', precision_score(testy, yhat_test1,
average='weighted'))
print ('\n classification report_test:\n', classification_report(testy, yhat_test1))
print ('\n confusion matrix_test:\n', confusion_matrix(testy, yhat_test1))

Accuracy_test: 0.7904761904761904
F1 score_test: 0.778577884745583
Recall_test: 0.7904761904761904
Precision_test: 0.8159020873306588

classification report_test:
      precision    recall  f1-score   support

   0      0.77      1.00      0.87       10
   1      1.00      1.00      1.00       15
   2      1.00      1.00      1.00        5
   3      0.56      1.00      0.71        5
   4      0.71      1.00      0.83        5
   5      1.00      0.40      0.57        5
   6      1.00      1.00      1.00        5
   7      0.38      0.56      0.45        9
   8      1.00      1.00      1.00        5
   9      1.00      1.00      1.00       15
  10      1.00      0.86      0.86        9
  11      0.86      0.86      0.86        7
  12      0.83      1.00      0.91        5
  13      1.00      0.57      0.73        7
  14      0.25      0.12      0.17        8

   accuracy      0.82      0.80      0.79      105
  macro avg      0.82      0.80      0.79      105
 weighted avg      0.82      0.79      0.78      105

confusion matrix_test:
[[10  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0 15  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  5  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  5  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  5  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  5  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  5  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0  5  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0  5  0  0  0  0  0  0  0]
 [  0  0  0  1  0  0  2  0  5  0  0  0  0  0  1]
 [  0  0  0  0  0  0  1  0  0  0  6  0  0  0  0]
 [  0  0  0  0  0  0  0  0  0  0  0  5  0  0  0]
 [  2  0  0  1  0  0  0  0  0  0  0  0  0  4  0]
 [  1  0  0  0  1  0  0  0  0  0  0  0  0  0  1]]
```

Here is a summary for both the **SVM** and **KNN** classifiers

		Metrics	Train	Test
KNN		Accuracy	1.0	0.79
		F1 score	0.99	0.77
		Recall	1.0	0.79
		Precision	1.0	0.81

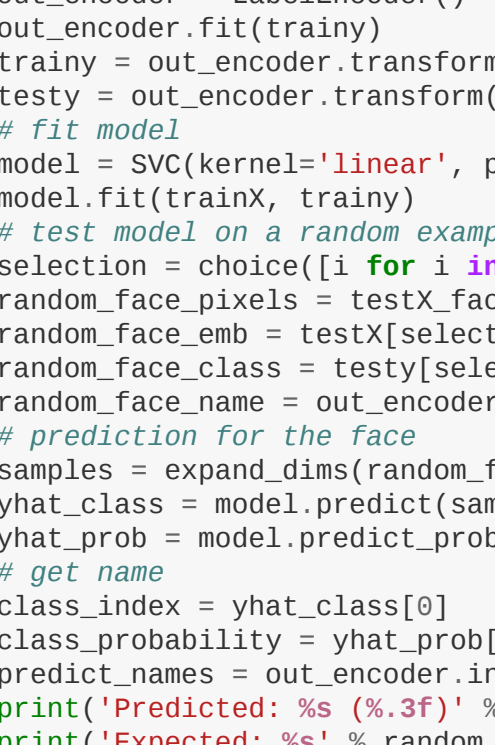
		Metrics	Train	Test
SVM		Accuracy	0.99	0.82
		F1 score	0.99	0.81
		Recall	0.99	0.82
		Precision	0.99	0.86

```
In [27]: # develop a classifier for the 5 Celebrity Faces Dataset
from random import choice
from numpy import choice
from numpy import expand_dims
from sklearn.preprocessing import LabelEncoder
from sklearn.svm import SVC
from matplotlib import pyplot

# load faces
data = load('masked-faces-dataset.npz')
testX_faces = data['arr_2']
# load face embeddings
data = load('masked-faces-embeddings.npz')
trainX, trainy, testX, testy = data['arr_0'], data['arr_1'], data['arr_2'], data['arr_3']
# normalize input vectors
in_encoder = Normalizer(norm='l2')
trainX = in_encoder.transform(trainX)
testX = in_encoder.transform(testX)
# label encode targets
out_encoder = LabelEncoder()
out_encoder.fit(trainy)
trainy = out_encoder.transform(trainy)
testy = out_encoder.transform(testy)
# fit model
model = SVC(kernel='linear', probability=True)
model.fit(trainX, trainy)
# test model on a random example from the test dataset
selection = choice([i for i in range(testX.shape[0])])
random_face_emb = testX_faces[selection]
random_face_name = testy[selection]
random_face_name = out_encoder.inverse_transform([random_face_name])
# prediction for the face
samples = expand_dims(random_face_emb, axis=0)
yhat_class = model.predict(samples)
yhat_prob = model.predict_proba(samples)
# get name
class_index = yhat_class[0]
class_probability = yhat_prob[0,class_index] * 100
predict_names = out_encoder.inverse_transform(yhat_class)
print('Predicted: %s (%.3f)' % (predict_names[0], class_probability))
print('Expected: %s' % random_face_name[0])

# plot for fun
pyplot.imshow(random_face_pixels)
title = '%s (%.3f)' % (predict_names[0], class_probability)
pyplot.title(title)
pyplot.show()

Predicted: dida (61.588)
Expected: dida
```



```
In [ ]:
```